

Process Management & Zombies

Benjamin Brewster

Except as noted, all images copyrighted with Creative Commons licenses,
with attributions given whenever available

Running Processes


- How can we tell which processes are running? Use the `ps` command to get information about currently running processes

- `Ps` by itself is really boring, and not all that useful:

```
$ ps
```

| PID | TTY | TIME | CMD |
|-------|-------|----------|------|
| 18779 | pts/8 | 00:00:00 | bash |
| 18934 | pts/8 | 00:00:00 | ps |

- I've put together my two favorite ways to run it on the next few slides



ps For Me

Reminder: these aliases go into your ~/.bashrc file

```
$ alias
```

```
...
```

```
alias psme='ps -o ppid,pid,euser,stat,%cpu,rss,args | head -n 1; ps -eH -o  
ppid,pid,euser,stat,%cpu,rss,args | grep brewsteb'
```

```
...
```

```
$ psme
```

| PPID | PID | EUSER | STAT | %CPU | RSS | COMMAND |
|-------|-------|----------|------|------|------|---|
| 4533 | 18776 | root | Ss | 0.2 | 4284 | sshd: brewsteb [priv] |
| 18776 | 18778 | brewsteb | S | 0.0 | 2112 | sshd: brewsteb@pts/8 |
| 18778 | 18779 | brewsteb | Ss | 0.0 | 2044 | -bash |
| 18779 | 18911 | brewsteb | R+ | 4.0 | 1840 | ps -eH -o ppid,pid,euser,stat,%cpu,rss,args |
| 18779 | 18912 | brewsteb | S+ | 0.0 | 820 | grep brewsteb |

- | | | | |
|---------|-------------------|-----------|---|
| • PPID | Parent Process ID | • %CPU | Percentage of CPU time this process occupies |
| • PID | Process ID | • RSS | Real Set Size - kilobytes of RAM in use by this process |
| • EUSER | Effective User ID | • Command | The actual command the user entered |
| • STAT | Execution State | | |

ps For Me

```
$ psme
```


| PPID | PID | EUSER | STAT | %CPU | RSS | COMMAND |
|-------|-------|----------|------|------|------|---|
| 4533 | 18776 | root | Ss | 0.2 | 4284 | sshd: brewsteb [priv] |
| 18776 | 18778 | brewsteb | S | 0.0 | 2112 | sshd: brewsteb@pts/8 |
| 18778 | 18779 | brewsteb | Ss | 0.0 | 2044 | -bash |
| 18779 | 18911 | brewsteb | R+ | 4.0 | 1840 | ps -eH -o ppid,pid,euser,stat,%cpu,rss,args |
| 18779 | 18912 | brewsteb | S+ | 0.0 | 820 | grep brewsteb |

First State Character:

- D Uninterruptible sleep (usually IO)
- R Running or runnable (on run queue)
- S Interruptible sleep (waiting for an event to complete)
- T Stopped, either by a job control signal or because it is being traced
- Z Defunct ("zombie") process, terminated but not reaped by its parent

Second State Character (Optional):

- < High-priority (not nice to other users)
- N Low-priority (nice to other users)
- L Has pages locked into memory (for real-time and custom IO)
- s Is a session leader (closes all child processes on termination)
- L Is multi-threaded (Uses `pthread`)
- + Is in the foreground process group



ps For All



```
$ alias
```

```
...
alias psall='ps -eH -o ppid,pid,euser,stat,%cpu,rss,args | awk '\''$1!=0'\'' | awk '\''$1!=1'\'' | awk
'\''$1!=2'\'' | more'
...
```

```
$ psall
```

| PPID | PID | EUSER | STAT | %CPU | RSS | COMMAND |
|-------|-------|----------|------|------|--------|---|
| ... | | | | | | |
| 4533 | 21922 | root | Ss | 0.0 | 4288 | sshd: meadosc [priv] |
| 21922 | 21936 | meadosc | S | 0.0 | 2128 | sshd: meadosc@pts/11 |
| 21936 | 21937 | meadosc | Ss | 0.0 | 2024 | -tcsh |
| 21937 | 21962 | meadosc | S+ | 0.0 | 1900 | bash |
| 4533 | 25083 | root | Ss | 0.4 | 4284 | sshd: brewsteb [priv] |
| 25083 | 25104 | brewsteb | S | 0.0 | 2112 | sshd: brewsteb@pts/8 |
| 25104 | 25105 | brewsteb | Ss | 0.0 | 2040 | -bash |
| 25105 | 25761 | brewsteb | R+ | 8.0 | 1852 | ps -eH -o ppid,pid,euser,stat,%cpu,rss,args |
| 25105 | 25762 | brewsteb | S+ | 0.0 | 908 | awk \$1!=0 |
| 25105 | 25763 | brewsteb | S+ | 0.0 | 908 | awk \$1!=1 |
| 25105 | 25764 | brewsteb | S+ | 0.0 | 908 | awk \$1!=2 |
| 25105 | 25765 | brewsteb | S+ | 0.0 | 708 | more |
| 4982 | 5194 | root | Ss | 0.0 | 5136 | /opt/dell/srvadmin/sbin/dsm_sa_datamgrd |
| 5339 | 5340 | root | Sl | 0.1 | 244404 | /opt/dell/srvadmin/sbin/dsm_om_connsvc -run |
| 5461 | 25756 | root | Zs | 0.0 | 0 | [check_nfs.sh] <defunct> |
| 23087 | 23088 | groveed | Ss+ | 0.0 | 1784 | -bin/tcsh |
| ... | | | | | | |


Zombie?

- When a child process terminates, but its parent does not wait for it, the process becomes known as a zombie (aka *defunct*)



Zombies!?!

- Child processes must report to their parents before their resources will be released by the OS
- If the parents aren't waiting for their children, the processes become the *living undead* – *forever consuming, forever enslaved to a non-life of waiting and watching*.
- The purpose of a zombie process is to retain the state that `wait()` can retrieve; they *want* to be harvested



Makin' Zombies

```
$ cat forkyouzombie.c
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>


void main()
{
    pid_t spawnPid = -5;
    int childExitStatus = -5;

    spawnPid = fork();
    switch (spawnPid)
    {
        case -1:
            perror("Hull Breach!\n");
            exit(1);
            break;

        case 0:
            printf("CHILD: Terminating!\n");
            break;

        default:
            printf("PARENT: making child a zombie for ten seconds;\n");
            printf("PARENT: Type \"ps -elf | grep 'username\\'\" to see the defunct child\n");
            printf("PARENT: Sleeping...\n");
            fflush(stdout);
            sleep(10);
            waitpid(spawnPid, &childExitStatus, 0);
            break;
    }
    printf("This will be executed by both of us!\n");
    exit(0);
}
```

Make sure all
text is outputted
before sleeping



Output - Makin' Zombies

```
$ gcc -o forkyouzombie forkyouzombie.c
```

```
$ forkyouzombie
```

```
PARENT: making child a zombie for ten seconds;
```

```
PARENT: Type "ps -elf | grep 'username'" to see the defunct child
```

```
PARENT: Sleeping...
```

```
CHILD: Terminating!
```

```
This will be executed by both of us!
```

```
This will be executed by both of us!
```

Dramatic ten-second pause right here...

```
// In a second terminal...
```


Nice, reasonable, short way to do a useful `ps`

```
$ ps -elf | grep 'brewsteb'
```

| | | | | | | | | | | | | | | |
|---|---|----------|-------|-------|----|----|---|---|-------|---------|-------|--------|----------|---------------------------|
| 4 | S | root | 15296 | 4443 | 0 | 80 | 0 | - | 32719 | unix_s | 10:55 | ? | 00:00:00 | sshd: brewsteb [priv] |
| 5 | S | brewsteb | 15298 | 15296 | 0 | 80 | 0 | - | 32719 | poll_s | 10:55 | ? | 00:00:00 | sshd: brewsteb@pts/40 |
| 0 | S | brewsteb | 15299 | 15298 | 0 | 80 | 0 | - | 30233 | wait | 10:55 | pts/40 | 00:00:00 | -bash |
| 0 | S | brewsteb | 17053 | 27991 | 0 | 80 | 0 | - | 981 | hrttime | 11:15 | pts/9 | 00:00:00 | forkyouzombie |
| 1 | Z | brewsteb | 17054 | 17053 | 0 | 80 | 0 | - | 0 | exit | 11:15 | pts/9 | 00:00:00 | [forkyouzombie] <defunct> |
| 0 | R | brewsteb | 17057 | 15299 | 12 | 80 | 0 | - | 30674 | - | 11:15 | pts/40 | 00:00:00 | ps -elf |
| 0 | S | brewsteb | 17058 | 15299 | 0 | 80 | 0 | - | 25829 | pipe_w | 11:15 | pts/40 | 00:00:00 | grep brewsteb |
| 4 | S | root | 27987 | 4443 | 0 | 80 | 0 | - | 32719 | unix_s | 08:51 | ? | 00:00:00 | sshd: brewsteb [priv] |
| 5 | S | brewsteb | 27990 | 27987 | 0 | 80 | 0 | - | 32719 | poll_s | 08:51 | ? | 00:00:00 | sshd: brewsteb@pts/9 |
| 0 | S | brewsteb | 27991 | 27990 | 0 | 80 | 0 | - | 30234 | wait | 08:51 | pts/9 | 00:00:00 | -bash |

How to Deal With Zombies

Zombies stay in the system
until they are waited for



How to deal with Zombies

Zombies stay in the system
until they are waited for




Orphan Zombies!

- If a parent process terminates *without* cleaning up its zombies, the zombies become orphan zombies
- Orphans are adopted by the `init` process (usually `pid = 1`) which periodically (in practice, very quickly) `wait()` for orphans
- Thus eventually, the orphan zombies die

kill

- This UNIX command is used to kill programs
 - another old version is called `kfork`
- “kill” is really a misnomer – it really *just sends signals*



kill



The PID of the process being signaled

```
kill -TERM 1234
```



The signal to send

- The given PID affects who the signal is sent to:
 - If $PID > 0$, then the signal will be sent to the process PID given
 - If $pid == 0$, then the signal is sent to all processes in the same process group as the sender (from an interactive command line, this means the foreground process group, i.e. your shell)
 - More trickiness for $pid < 0$
- We'll discuss more signals later, but you can use the signal `KILL` to tell a process to immediately terminate with no clean-up

top


- `top` allows you to view the processes running on the machine in real time - one of the few animated built-in programs

\$ top

```
top - 14:14:34 up 34 days,  5:15,  9 users,  load average: 0.03, 0.18, 0.22
Tasks: 703 total,   1 running, 697 sleeping,   4 stopped,   1 zombie
Cpu(s):  0.1%us,   0.1%sy,   0.0%ni, 99.8%id,   0.0%wa,   0.0%hi,   0.0%si,   0.0%st
Mem:  65922540k total,  7876576k used, 58045964k free,   663988k buffers
Swap:  2588668k total,           0k used,  2588668k free,  5258716k cached
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-------|----------|----|----|-------|------|------|---|------|------|---------|-------------|
| 27609 | brewsteb | 20 | 0 | 27884 | 1796 | 996 | R | 3.4 | 0.0 | 0:00.10 | top |
| 1 | root | 20 | 0 | 33656 | 1624 | 1292 | S | 0.0 | 0.0 | 1:41.47 | init |
| 2 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:09.18 | kthreadd |
| 3 | root | RT | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:21.23 | migration/0 |
| 4 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:14.47 | ksoftirqd/0 |
| 5 | root | RT | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00.00 | stopper/0 |

...




Diagnosing a Slow CPU

- The *uptime* command shows the average number of runnable processes over several different periods of time (the same info top displays)

```
$ uptime
```

```
1:23pm up 25 day(s), 5:59, 72 users, load average: 0.18, 0.19, 0.20
```

- This shows the average number of runnable (the current running process plus the queue of processes waiting to be run) or uninterruptable (waiting for IO) processes over the last 1, 5 and 15 minutes
- If uptime is showing that your runnable queue is consistently *larger than the number of cores*, your CPU is a bottleneck and is causing slow-down



Diagnosing a Slow CPU - Number of Cores?

```
$ cat /proc/cpuinfo
```

```
processor      : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 45
model name    : Intel(R) Xeon(R) CPU E5-2665 0 @ 2.40GHz
stepping      : 7
microcode     : 1808
cpu MHz       : 2399.993
cache size    : 20480 KB
physical id   : 0
siblings      : 16
core id       : 0
cpu cores    : 8
apicid        : 0
initial apicid : 0
fpu           : yes
fpu_exception : yes
```

```
...
```



Diagnosing a Slow CPU - Example - Single Core

```
$ uptime
```

```
14:33:04 up 34 days,  5:34, 10 users,  load average: 0.05, 0.15, 0.20
```

This CPU is the champ... or it's not being given anything to do

```
$ uptime
```

```
14:33:04 up 34 days,  5:34, 10 users,  load average: 0.88, 1.03, 0.96
```

This CPU is at max - time to upgrade!

```
$ uptime
```

```
14:33:04 up 34 days,  5:34, 10 users,  load average: 4.79, 7.23, 6.44
```

It's 3am, and your server is borked; start paging everyone!

Diagnosing a Slow CPU - Example - Octo Core

```
$ uptime
```

```
14:33:04 up 34 days, 5:34, 10 users, load average: 0.05, 0.15, 0.20
```

Your CPU is bored, and you wasted all the money; but hey, headroom for games!

```
$ uptime
```

```
14:33:04 up 34 days, 5:34, 10 users, load average: 7.99, 8.10, 7.94
```

This CPU is handling processes exactly as fast as it gets them - time to make it more betterer


```
$ uptime
```

```
14:33:04 up 34 days, 5:34, 10 users, load average: 39.90, 41.54, 40.72
```

You're so fired

Job Control

- How do we start a program, and *still retain access* to the command line for the next program we want to run?
- Can we run multiple processes at once?
- This is called Job Control in UNIX-speak



Foreground/Background

- There can be only one shell ***foreground*** process – it's the one you're currently interacting with
- If you're at the command prompt, then your foreground program is the shell itself!

- Processes in the ***background*** can still be executing, but they can also be in any number of **stopped states**:

First State Character:

- D Uninterruptible sleep (usually IO)
- R Running or runnable (on run queue)
- S Interruptible sleep (waiting for an event to complete)
- T Stopped, either by a job control signal or because it is being traced
- Z Defunct ("zombie") process, terminated but not reaped by its parent

Foreground/Background in Reality


- There really isn't any difference between processes in these two states; its merely shell nomenclature used to distinguish between them
- When a user enters a command that is intended to run in the foreground (i.e. a normal command), the process started runs to completion *before* the user is prompted again
- When a user enters a command that is intended to run in the background (see later slides), the user is *immediately* prompted again after the process is executed
- In other words, control input to the terminal is not interrupted by a background process

Start Backgrounded

- Here's how to start a program in the background in the first place:

```
$ ping www.oregonstate.edu &
```

- The ampersand means to start in the background, and must be the last character
- Note that stdout and stderr are still going to the terminal for that process, and stdin might be too if the shell is badly programmed



Stopping a Process

- Sending the TSTP signal stops (not terminates) a process, and puts it into the background
 - Control-z also sends this signal

ping is now the foreground process, and I can't enter commands anymore

```
$ ping www.oregonstate.edu
```

```
PING www.orst.edu (128.193.4.112) 56(84) bytes of data.
```

```
64 bytes from www.orst.edu (128.193.4.112): icmp_seq=1 ttl=250 time=0.362 ms
```

```
64 bytes from www.orst.edu (128.193.4.112): icmp_seq=2 ttl=250 time=0.321 ms
```

```
64 bytes from www.orst.edu (128.193.4.112): icmp_seq=3 ttl=250 time=0.324 ms
```

```
64 bytes from www.orst.edu (128.193.4.112): icmp_seq=4 ttl=250 time=0.328 ms
```

```
^Z
```

```
[1]+  Stopped
```

```
ping www.oregonstate.edu
```

```
$
```

Our shell is once again the foreground process

jobs

- Use the `jobs` command to see what you're running:

```
$ ping www.oregonstate.edu
PING www.orst.edu (128.193.4.112) 56(84) bytes of data.
64 bytes from www.orst.edu (128.193.4.112): ...
64 bytes from www.orst.edu (128.193.4.112): ...
^Z
[3]+  Stopped                  ping www.oregonstate.edu

$ jobs -l
[1]-  31314 Stopped              ping www.oregonstate.edu
[2]-  31317 Stopped              ping www.oregonstate.edu
[3]+  31327 Stopped              ping www.oregonstate.edu

$ kill -KILL 31327
[3]+  Killed                    ping www.oregonstate.edu

$ kill -KILL %1
[1]-  Killed                    ping www.oregonstate.edu
```

The `-l` switch adds the PID

The `-` symbol means it was the second-to-last process put in the background

The `+` symbol means it was the last process put in the background

"Job 1"

`fg`

- Use the job numbers provided by `jobs` to manipulate processes
- Bring job 1 from the background to the foreground, and start it running again:

`fg %1`

- Bring most recent backgrounded job to the foreground, and start it running again:

`fg`



bg

- Start a specific stopped program that is currently in the background (and keep it in the background):

```
bg %1
```

- Start the most recently stopped program in the background (and keep it in the background):

```
bg
```




Who's Got Control of stdout?

- Be advised – background processes can still write to any file including stdout & stderr!

- Jobs demo:

```
1. ping www.oregonstate.edu
2. CTRL-Z
3. jobs
4. fg %1
5. CTRL-Z
6. jobs
7. bg %1
8. ps
9. CTRL-Z CTRL-Z CTRL-Z (doesn't do anything)
10. fg %1
11. CTRL-C
```



You're Suspended

- Suspend a process that is currently running in the background when you're at the shell

Send stderr and stdout somewhere other than the terminal

Background this command!

```
$ ping www.oregonstate.edu 2>/dev/null 1>logfile &
```

```
[1] 1660
```

```
$ jobs
```

```
[1]+  Running
```

```
ping www.oregonstate.edu 2> /dev/null > /dev/null &
```

```
$ kill -TSTP %1
```

```
[1]+  Stopped
```

```
ping www.oregonstate.edu 2> /dev/null > /dev/null
```

```
$ jobs
```


```
[1]+  Stopped
```

```
ping www.oregonstate.edu 2> /dev/null > /dev/null
```

history - a Command Visibility Utility

- The history command provides a listing of your previous commands:

```
$ history 5  
1012  jobs  
1013  psme  
1014  top  
1015  jobs  
1016  history 5
```



Execute a Previous Command

```
$ history 3
1030 jobs
1031 psme
1032 history 3
$ !1030
jobs
$ history 3
1032 history 3
1033 jobs
1034 history 3
$ !-2
jobs
$ !!
jobs
$ history 3
1034 history 3
1035 jobs
1036 history 3
```

Note that no exclamation marks are in the history list - only the actual commands, even when repeated with the ! operator