


More UNIX I/O

Benjamin Brewster

Except as noted, all images copyrighted with Creative Commons licenses,
with attributions given whenever available


Open File Inheritance

- When you `exec...()` a process, you replace the current process with a new one, but the files are still open and accessible to the new process
- This may not be what you want!
- In this lecture, we discuss other IPC methods and UNIX I/O. We'll cover:
 - Close On Exec
 - File redirection in C
 - Pipes




I/O redirection

- We saw I/O redirection in the shell:
 - `ls > file`
 - `stats < file1`
 - `cat longfile | more`
 - `find . -name "paper" -print 2> /dev/null`
 - `echo "an error occurred" 1>&2`
- I/O redirection is possible *because* open files are shared across `fork()` and `exec...()`



I/O redirection

- I/O redirection is possible *because* open files are shared across `fork()` and `exec...()`
- Each child process has the same files open as its parent, as the file descriptors are the same file descriptors in both
- Note this important point: both parent and child read & write function calls move the *same* file pointer for the shared file descriptor!
 - To prevent this, have one of the processes close and re-open the file



Sharing File Pointers - Example 1 of 2

```
$ cat forkFPsharing.c
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>

void main()
{
    pid_t forkPID;
    int childExitMethod;
    int fileDescriptor;
    char *newFilePath = "./newFile.txt";
    char readBuffer[8];
    memset(readBuffer, '\0', sizeof(readBuffer));

    printf("PARENT: Opening file.\n");
    fileDescriptor = open(newFilePath, O_RDWR | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
    if (fileDescriptor == -1) { printf("Hull breach - open() failed on \"%s\"\n", newFilePath); exit(1); }

    printf("PARENT: Writing 01234 to file.\n");
    write(fileDescriptor, "01234", 5);
    printf("PARENT: New FP position (all FP positions are zero-indexed): %d\n", lseek(fileDescriptor, 0, SEEK_CUR));
    fflush(stdout);

    printf("PARENT: Spawning child.\n");
    forkPID = fork();
```

The child process will have `fileDescriptor` open, with the SHARED file pointer!

Sharing File Pointers - Example 2 of 2

```
switch (forkPID)
{
case -1: perror("Hull Breach!"); exit(1); break;
case 0:
    printf("CHILD: Writing C to file.\n"); fflush(stdout);
    write(fileDescriptor, "C", 1);
    printf("CHILD: After write, new FP position: %d\n", lseek(fileDescriptor, 0, SEEK_CUR)); fflush(stdout);
    printf("CHILD: lseek back 3 chars.\n"); fflush(stdout);
    printf("CHILD: New FP position: %d\n", lseek(fileDescriptor, -3, SEEK_CUR)); fflush(stdout);
    printf("CHILD: Reading char.\n"); fflush(stdout);
    read(fileDescriptor, &readBuffer, 1);
    printf("CHILD: After read, new FP position: %d, char read was: %c\n", lseek(fileDescriptor, 0, SEEK_CUR), readBuffer[0]); fflush(stdout);
    break;
default:
    printf("PARENT: Writing P to file.\n"); fflush(stdout);
    write(fileDescriptor, "P", 1);
    printf("PARENT: After write, new FP position: %d\n", lseek(fileDescriptor, 0, SEEK_CUR)); fflush(stdout);
    printf("PARENT: lseek back 3 chars.\n"); fflush(stdout);
    printf("PARENT: New FP position: %d\n", lseek(fileDescriptor, -3, SEEK_CUR)); fflush(stdout);
    printf("PARENT: Reading char.\n"); fflush(stdout);
    read(fileDescriptor, &readBuffer, 1);
    printf("PARENT: After read, new FP position: %d, char read was: %c\n", lseek(fileDescriptor, 0, SEEK_CUR), readBuffer[0]); fflush(stdout);
    waitpid(forkPID, &childExitMethod, 0);
    lseek(fileDescriptor, 0, SEEK_SET);
    read(fileDescriptor, &readBuffer, 7);
    printf("PARENT: child terminated; file contents: %s\n", readBuffer); fflush(stdout);
    break;
}
```

This code and trace statements just reads and writes to a file open and shared in both the parent and child processes

Sharing File Pointers - Results

```
$ gcc -o forkFPsharing forkFPsharing.c
```

```
$ forkFPsharing
```

```
PARENT: Opening file.
```

```
PARENT: Writing 01234 to file.
```

```
PARENT: New FP position (all FP positions are zero-indexed): 5
```

```
PARENT: Spawning child.
```

```
PARENT: Writing P to file.
```

```
PARENT: After write, new FP position: 6
```

```
PARENT: lseek back 3 chars.
```

```
PARENT: New FP position: 3
```

```
PARENT: Reading char.
```

```
PARENT: After read, new FP position: 4, char read was: 3
```

```
CHILD: Writing C to file.
```

```
CHILD: After write, new FP position: 5
```


```
CHILD: lseek back 3 chars.
```

```
CHILD: New FP position: 2
```

```
CHILD: Reading char.
```


```
CHILD: After read, new FP position: 3, char read was: 2
```

```
PARENT: child terminated; file contents: 0123CP
```




Important Background Review

- The kernel opens `stdin`, `stdout`, and `stderr` automatically for every process created:
 - File descriptor `0` is `stdin`
 - File descriptor `1` is `stdout`
 - File descriptor `2` is `stderr`
- They default to reading and writing to the terminal
- The trick: you can change where the standard I/O streams are pointing to at any time before the `exec...()` call, including after the `fork()`




Redirecting stdout

How all processes start:




Redirecting stdout

1. First open the new file



Redirecting stdout

2. Call `dup2 ()` to change fd 1 to point where fd 3 points: `dup2 (3, 1) ;`



stdout, merely another name for fd 1, can now be used to access "myfile.data"

Redirecting stdout

Set FD 1 (stdout) to point to the same place that targetFD points.

So, anything written to stdout (like with `printf()`) will go to targetFD, which is the passed in filename

```
$ cat redirectToFile.c
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>

int main(int argc, char* argv[])
{
    if (argc == 1)
    {
        printf("Usage: redirectToFile <filename to redirect stdout to>\n");
        exit(1);
    }

    int targetFD = open(argv[1], O_WRONLY | O_CREAT | O_TRUNC, 0644);
    if (targetFD == -1) { perror("open()"); exit(1); }
    printf("targetFD == %d\n", targetFD); // Written to terminal

    int result = dup2(targetFD, 1);
    if (result == -1) { perror("dup2"); exit(2); }
    printf("targetFD == %d, result == %d\n", targetFD, result); // Written to file

    return(0);
}

$ gcc -o redirectToFile redirectToFile.c
$ redirectToFile
Usage: redirectToFile <filename to redirect stdout to>
$ redirectToFile test.junk
targetFD == 3
$ cat test.junk
targetFD == 3, result == 1
```

Redirecting stdout & stdin with `execvp()`

```
$ cat sortViaFiles.c
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
```

```
int main(int argc, char* argv[])
{
    int sourceFD, targetFD, result;

    if (argc != 3)
    {
        printf("Usage: sortViaFiles <input filename> <output filename>\n");
        exit(1);
    }

    sourceFD = open(argv[1], O_RDONLY);
    if (sourceFD == -1) { perror("source open()"); exit(1); }
    printf("sourceFD == %d\n", sourceFD); // Written to terminal

    targetFD = open(argv[2], O_WRONLY | O_CREAT | O_TRUNC, 0644);
    if (targetFD == -1) { perror("target open()"); exit(1); }
    printf("targetFD == %d\n", targetFD); // Written to terminal

    result = dup2(sourceFD, 0);
    if (result == -1) { perror("source dup2()"); exit(2); }
    result = dup2(targetFD, 1);
    if (result == -1) { perror("target dup2()"); exit(2); }

    execvp("sort", "sort", NULL);
    return(3);
}
```

Set FD 0 (stdin) to point to the same place that `sourceFD` points.

So, anything in stdin (like the contents of the input file) will be used by this program

Set FD 1 (stdout) to point to the same place that `targetFD` points.

So, anything written to stdout (like the result of sort) will go to `targetFD`, which is the passed in output filename

`execvp()` now starts with stdin and stdout pointing to files, which are used by `sort`

Redirecting stdout & stdin with `execvp()`

```
$ cat sortViaFiles.c
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
```

```
int r
{
```

Results:

```
$ gcc -o sortViaFiles sortViaFiles.c
```

```
$ echo -e "3\n1\n2" > junkinput
```

```
$ cat junkinput
```

```
3
```

```
1
```

```
2
```

```
$ sortViaFiles junkinput junkoutput
```

```
sourceFD == 3
```

```
targetFD == 4
```

```
$ cat junkoutput
```

```
1
```

```
2
```

```
3
```

```
execvp("sort", "sort", NULL);
return(3);
```

```
}
```

Set FD 0 (stdin) to point to the same place that `sourceFD` points.

So, anything in stdin (like the contents of the input file) will be used by this program

Set FD 1 (stdout) to point to the same place that `targetFD` points.

So, anything written to stdout (like the result of sort) will go to `targetFD`, which is the passed in output filename

```
name> <output filename>\n");
```

```
exit(1); }
```

```
h to terminal
```

```
_TRUNC, 0644);
```

```
exit(1); }
```

```
h to terminal
```

```
it(2); }
```

```
it(2); }
```

`execvp()` now starts with stdin and stdout pointing to files, which are used by `sort`

Close On Exec Details

- A flag specified in an `open()` call that tells the kernel to close any open files when/if the process gets `exec . . . ()`'d
- Why do we care? Because open files are inherited by child processes
 - Thus the file pointer is shared (where the file is currently being read and written to)
 - Secure and/or sensitive data is shared with the new process
- Specific to only that file you added the flag to
- This is inherited through `fork`, so if the parent specifies it, a child that later calls `exec...()` will trigger the close on exec flag

Close on exec example

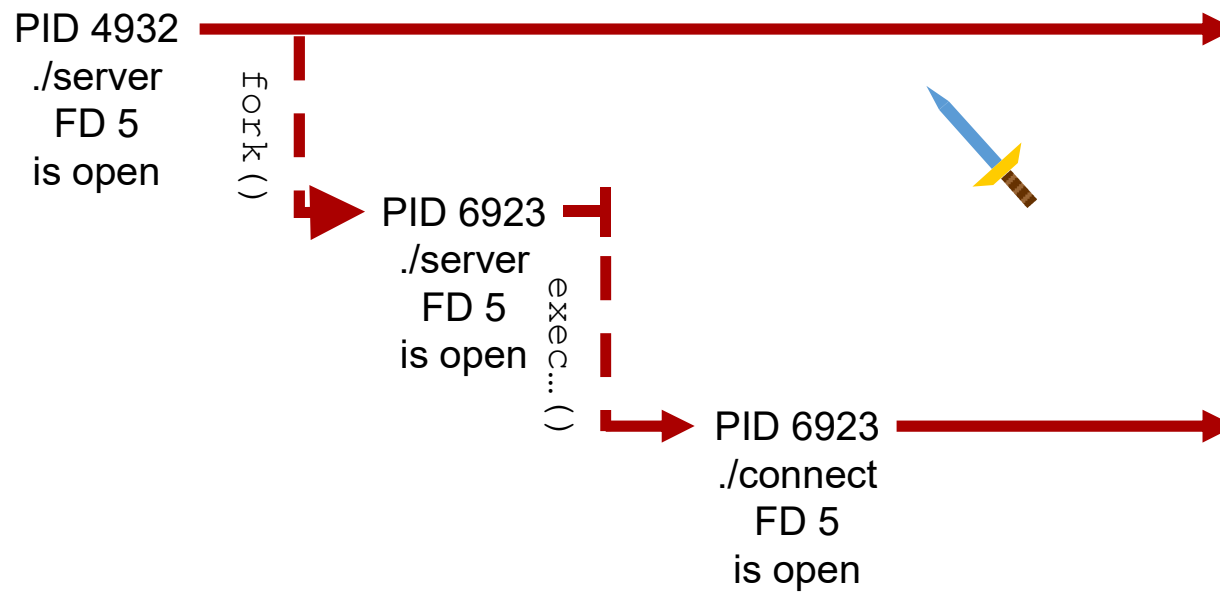
```
#include <fcntl.h>

...
int fd;
fd = open("file", O_RDONLY);
...
fcntl(fd, F_SETFD, FD_CLOEXEC);
...
// exec...
```

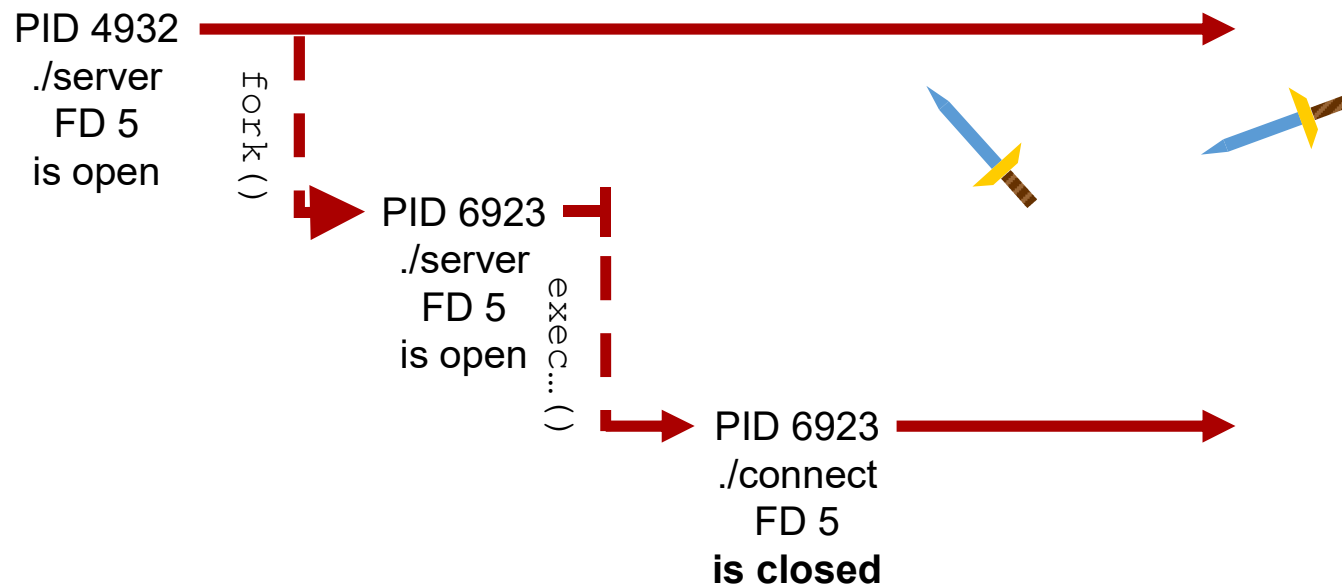


Close on Exec

Normally...




With Close On Exec




Real Inter-Process Communication (IPC)

- IPC methods in UNIX
 - **Intermediate/temporary files** :: Often used together with I/O redirection
 - **Pipes** :: IPC between two processes forked by a common ancestor process
 - **FIFOs** (named pipes) :: communication between any two processes on the same machine
 - **Message queues** :: communication between any two processes on the same machine
 - Not common; older System V libraries were replaced with POSIX libraries
 - Not a simple byte stream: In Linux, the POSIX queues can be mounted as a filesystem, with each message appearing as a file; can then use `ls`, `rm`, etc.
 - Supports message categories - often used for priorities
 - **Sockets** :: communication between any two processes, potentially separated by a network




Between-Process IPC - Intro to Pipes

- I/O redirection with `dup2 ()` allows you to redirect input and output between processes and files...
- But how do we *redirect input* between processes and other processes on the same machine?
 - We could use temporary/intermediate files, but:
 - Writes to disk are slow
 - No fast & efficient way to track when the other process is ready to receive or send new data other than some sort of semaphore library
 - Better answer: use **pipes**!



Pipes


- Pipes provide a way to connect a write-only file descriptor in one process to a read-only file descriptor in another process



- `write()` puts bytes in the pipe, `read()` takes them out

Creating a Pipe

- Pipes are possible because file descriptors are shared across `fork()` and `exec...()`
- A parent process creates a pipe
 - Results in two new open file descriptors, one for input and one for output
- The parent process calls `fork()` and possibly `exec...()`
 - Parent and child have the file descriptors created with the pipe
- The child process now reads from the input file descriptor, and the parent process writes to the output file descriptor
 - or vice-versa



The `pipe()` Function


- You pass `pipe()` an array of two integers, where it stores the two new open file descriptors that it creates
- The first is the input file descriptor, and the second is the output file descriptor
- One of the descriptors should be used by the parent process and the other should be used by the child process



We'll talk about how to use a pipe to communicate between two non-decent processes later


Flow Control with `read()`

- `read()` succeeds if data is available
 - Receives the data and returns immediately
 - The return value of `read()` tells you how many bytes were read, which *may be less than you requested*
- If data is not available, `read()` will *block* waiting for data (your process execution is suspended until data arrives)
 - `read()` is a system call



Flow Control with `write()`

- Similarly, `write` will not return until all the data has been written
 - `write()` is a system call
- Pipes have a certain size
 - Only so much data will fit in a pipe (typically 64K, but can be changed)
 - If the pipe fills up, and there is no more room, `write()` will *block* until space becomes available (ie somebody `reads` the data from the pipe)



pipe () Example

```
$ cat pipeNfork.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <string.h>
void main()
{
    int r, pipeFDs[2];
    char completeMessage[512], readBuffer[10];
    pid_t spawnpid;

    if (pipe(pipeFDs) == -1) { perror("Hull Breach!"); exit(1); } // Create the pipe with error check

    spawnpid = fork(); // Fork the child, which will write into the pipe
    switch (spawnpid)
    {
    case 0: // Child
        close(pipeFDs[0]); // close the input file descriptor
        write(pipeFDs[1], "CHILD: Hi parent!@", 19); // Write the entire string into the pipe
        exit(0); break; // Terminate the child
    default: // Parent
        close(pipeFDs[1]); // close output file descriptor
        memset(completeMessage, '\0', sizeof(completeMessage)); // Clear the buffer

        while (strstr(completeMessage, "@@") == NULL) // As long as we haven't found the terminal...
        {
            memset(readBuffer, '\0', sizeof(readBuffer)); // Clear the buffer
            r = read(pipeFDs[0], readBuffer, sizeof(readBuffer) - 1); // Get the next chunk
            strcat(completeMessage, readBuffer); // Add that chunk to what we have so far
            printf("PARENT: Message received from child: \"%s\", total: \"%s\\n\", readBuffer, completeMessage);
            if (r == -1) { printf("r == -1\\n"); break; } // Check for errors
            if (r == 0) { printf("r == 0\\n"); break; }
        }
        int terminalLocation = strstr(completeMessage, "@@") - completeMessage; // Where is the terminal
        completeMessage[terminalLocation] = '\0'; // End the string early to wipe out the terminal
        printf("PARENT: Complete string: \"%s\\n\", completeMessage);
        break;
    }
}

$ pipeNfork
PARENT: Message received from child: "CHILD: Hi", total: "CHILD: Hi"
PARENT: Message received from child: " parent!@", total: "CHILD: Hi parent!@"
PARENT: Message received from child: "@", total: "CHILD: Hi parent!@"
PARENT: Complete string: "CHILD: Hi parent!"
```

```
$ cat pipeNfork.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <string.h>
void
{
```

Pointer Arithmetic Example

```
int array[3];
```

Element	Address Calculation	Address
array	1100	1100
array[0]	1100 + (sizeof(int) * 0)	1100
array[1]	1100 + (sizeof(int) * 1)	1104
array[2]	1100 + (sizeof(int) * 2)	1108

Therefore:

```
int* x = &array[0];           // x is int pointer
int* y = &array[2];           // y is int pointer
int z = y - x == 1108 - 1100 == 8 (bytes) => 2 (ints); // z is int; z = 2
```

```

    strcat(completeMessage, readBuffer); // Add
    printf("PARENT: Message received from child: %s\n", readBuffer, completeMessage);
    if (r == -1) { printf("r == -1\n"); break; }
    if (r == 0) { printf("r == 0\n"); break; }
}
int terminalLocation = strstr(completeMessage, "@@") - completeMessage; // Where is the terminal
completeMessage[terminalLocation] = '\0'; // End the string early to wipe out the terminal
printf("PARENT: Complete string: \"%s\"\n", completeMessage);
break;
}
}

```

(4068 - 4000) / sizeof(int) = 17

```
$ pipeNfork
```

```
PARENT: Message received from child: "CHILD: Hi", total: "CHILD: Hi"
PARENT: Message received from child: " parent!@", total: "CHILD: Hi parent!@"
PARENT: Message received from child: "@", total: "CHILD: Hi parent!@"
PARENT: Complete string: "CHILD: Hi parent!"
```

```
$ cat pipeNfork.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <string.h>
void
{
```

Could also have just done this:

```
char* t = strstr(completeMessage, "@@"); // Where is the terminal
*t = '\\0';
```

```

        strcat(completeMessage, readBuffer); // Add
        printf("PARENT: Message received from child: %s\n", readBuffer, completeMessage);
        if (r == -1) { printf("r == -1\n"); break; }
        if (r == 0) { printf("r == 0\n"); break; }
    }
    int terminalLocation = strstr(completeMessage, "@@") - completeMessage; // Where is the terminal
    completeMessage[terminalLocation] = '\\0'; // End the string early to wipe out the terminal
    printf("PARENT: Complete string: \"%s\\n\\n\", completeMessage);
    break;
}
}


```

\$ pipeNfork

```
PARENT: Message received from child: "CHILD: Hi", total: "CHILD: Hi"
PARENT: Message received from child: " parent!@", total: "CHILD: Hi parent!@"
PARENT: Message received from child: "@", total: "CHILD: Hi parent!@"
PARENT: Complete string: "CHILD: Hi parent!"
```


Error Checking Reads and Writes

- Checking the return value of `read()` is very important
 - Not just if return value is -1 (an error)
 - The return value will tell you if the desired number of bytes was not read - this can tell you if the pipe didn't have the amount of data you expected it to
 - Our previous example used a terminator @@ instead of tracking byte counts because often you don't know how many bytes there will be
- Same goes for `write()`
 - If the number of bytes returned isn't what you expected (for example, if a signal handler interrupted), you'll need to loop over the write again, *writing only what got missed again* to it




Closing Pipes

- Process 1 closes output pipe:
 - If process 2 is currently blocked on a `read()`, then process 2's `read()` will return 0
- Process 2 closes input pipe:
 - If process 1 tries to write to the pipe, `write()` will return -1, and `errno` (in process 1) will be set to `EPIPE`
 - Process 1 will be sent the `SIGPIPE` signal




Named Pipe = FIFO

- FIFO = First-In, First-out
- Essentially, a persistent pipe, which is represented by a special file
- Create in C with `mkfifo()`, or with `mkfifo` in bash
- Once created, any process can open a FIFO with `open()`
- Once opened, it works just like a pipe (or really: just like any file)



FIFO Use Cases

- You want to build a client-server architecture on a single machine, but you don't want to deal with the complexities of sockets
 - Can be used to transmit data between two non-related processes that didn't use `pipe()` and then `fork()`
- You want to transmit data with a non-network aware program




FIFO shell example

```
$ mkfifo my_fifo
```

```
$ ls -l *my*
```


```
prw-rw----. 1 brewsteb upg57541 0 Oct 31 14:46 my_fifo
```

- Since they are files, you can apply most of the common bash shell commands like:
 - read, sort, wc, cut, awk, etc.
- As well as all of the common file input/output system calls in C: `open()`, `read()`, `write()`, etc.



Opening a FIFO

- When opening a FIFO, `open()` called by the first process will block; the first process will unblock once the second process calls `open`
- Example:
 1. Process A calls `open(..., O_RDONLY)` // Process A blocks
 2. Process B calls `open(..., O_WRONLY)` // Process A & B continue
// execution



FIFO Example

```
$ cat pipeNforkFIFO.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <string.h>
#include <fcntl.h>
void main()
{
    int r, newfifo, fd;
    char completeMessage[512], readBuffer[10];
    char stringToWrite[20] = "CHILD: Hi parent!@@";
    pid_t spawnpid;
    char* FIFOfilename = "myNewFifo";
    newfifo = mkfifo(FIFOfilename, 0644); // Create the FIFO
    spawnpid = fork(); // Fork the child, which will write into the pipe
    switch (spawnpid) {
        case 0: // Child
            fd = open(FIFOfilename, O_WRONLY); // Open the FIFO for writing
            if (fd == -1) { perror("CHILD: open()"); exit(1); }
            write(fd, stringToWrite, strlen(stringToWrite)); // Write the entire string into the pipe
            exit(0); break; // Terminate the child
        default: // Parent
            fd = open(FIFOfilename, O_RDONLY); // Open the FIFO for reading
            if (fd == -1) { perror("PARENT: open()"); exit(1); }
            memset(completeMessage, '\0', sizeof(completeMessage)); // Clear the buffer
            while (strstr(completeMessage, "@@") == NULL) { // As long as we haven't found the terminal...
                memset(readBuffer, '\0', sizeof(readBuffer)); // Clear the buffer
                r = read(fd, readBuffer, sizeof(readBuffer) - 1); // Get the next chunk
                strcat(completeMessage, readBuffer); // Add that chunk to what we have so far
                printf("PARENT: Message received from child: \"%s\\\", total: \"%s\\\"\\n\", readBuffer, completeMessage);
                if (r == -1) { printf("PARENT: r == -1, exiting\\n"); break; } // Check for errors
                if (r == 0) { printf("PARENT: r == 0, exiting\\n"); break; }
            }
            int terminalLocation = strstr(completeMessage, "@@") - completeMessage; // Where is the terminal
            completeMessage[terminalLocation] = '\0'; // End the string early to wipe out the terminal
            printf("PARENT: Complete string: \"%s\\\"\\n\", completeMessage);
            remove(FIFOfilename); // Delete the FIFO
            break;
    }
}

$ pipeNforkFIFO
PARENT: Message received from child: "CHILD: Hi", total: "CHILD: Hi"
PARENT: Message received from child: " parent!@", total: "CHILD: Hi parent!@"
PARENT: Message received from child: "@", total: "CHILD: Hi parent!@"
PARENT: Complete string: "CHILD: Hi parent!"
```